

## Importing the Dependencies

```
import numpy as np
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score
```

## Data Collection and Processing

```
# loading the csv data to a Pandas DataFrame
heart_data = pd.read_csv('/content/heart_disease_data.csv')
```

```
# print first 5 rows of the dataset
heart_data.head()
```

```

age  sex  cp  trestbps  chol  fbs  restecg  thalach  exang  oldpeak  slope  ca  thal  target
0    63   1   3     145   233    1         0     150     0        2.3    0  0    1      1
1    37   1   2     130   250    0         1     187     0        3.5    0  0    2      1
2    41   0   1     130   204    0         0     172     0        1.4    2  0    2      1
3    56   1   1     120   236    0         1     178     0        0.8    2  0    2      1
4    57   0   0     120   354    0         1     163     1        0.6    2  0    2      1
```

```
# print last 5 rows of the dataset
heart_data.tail()
```

```

age  sex  cp  trestbps  chol  fbs  restecg  thalach  exang  oldpeak  slope  ca  thal  target
298  57   0   0     140   241    0         1     123     1        0.2    1  0    3      0
299  45   1   3     110   264    0         1     132     0        1.2    1  0    3      0
300  68   1   0     144   193    1         1     141     0        3.4    1  2    3      0
301  57   1   0     130   131    0         1     115     1        1.2    1  1    3      0
302  57   0   1     130   236    0         0     174     0        0.0    1  1    2      0
```

```
# number of rows and columns in the dataset
heart_data.shape
```

```
(303, 14)
```

```
# getting some info about the data
heart_data.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 303 entries, 0 to 302
Data columns (total 14 columns):
 #   Column        Non-Null Count  Dtype  
---  -
 0   age           303 non-null    int64  
 1   sex           303 non-null    int64  
 2   cp            303 non-null    int64  
 3   trestbps      303 non-null    int64  
 4   chol          303 non-null    int64  
 5   fbs           303 non-null    int64  
 6   restecg       303 non-null    int64  
 7   thalach       303 non-null    int64  
 8   exang         303 non-null    int64  
 9   oldpeak       303 non-null    float64 
10  slope         303 non-null    int64  
11  ca            303 non-null    int64  
12  thal          303 non-null    int64  
13  target        303 non-null    int64  
dtypes: float64(1), int64(13)
memory usage: 33.3 KB
```

+ Code

+ Text



McAfee | WebAdvisor



Your download's being scanned.  
We'll let you know if there's an issue.

```
# checking for missing values
heart_data.isnull().sum()
```

	0
age	0
sex	0
cp	0
trestbps	0
chol	0
fbs	0
restecg	0
thalach	0
exang	0
oldpeak	0
slope	0
ca	0
thal	0
target	0

```
# statistical measures about the data
heart_data.describe()
```

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak	slope	
count	303.000000	303.000000	303.000000	303.000000	303.000000	303.000000	303.000000	303.000000	303.000000	303.000000	303.000000	303.000000
mean	54.366337	0.683168	0.966997	131.623762	246.264026	0.148515	0.528053	149.646865	0.326733	1.039604	1.399340	
std	9.082101	0.466011	1.032052	17.538143	51.830751	0.356198	0.525860	22.905161	0.469794	1.161075	0.616226	
min	29.000000	0.000000	0.000000	94.000000	126.000000	0.000000	0.000000	71.000000	0.000000	0.000000	0.000000	
25%	47.500000	0.000000	0.000000	120.000000	211.000000	0.000000	0.000000	133.500000	0.000000	0.000000	1.000000	
50%	55.000000	1.000000	1.000000	130.000000	240.000000	0.000000	1.000000	153.000000	0.000000	0.800000	1.000000	
75%	61.000000	1.000000	2.000000	140.000000	274.500000	0.000000	1.000000	166.000000	1.000000	1.600000	2.000000	
max	77.000000	1.000000	3.000000	200.000000	564.000000	1.000000	2.000000	202.000000	1.000000	6.200000	2.000000	

```
# checking the distribution of Target Variable
heart_data['target'].value_counts()
```

	count
target	
1	165
0	138

Splitting the Features and Target

```
X = heart_data.drop(columns='target', axis=1)
Y = heart_data['target']
```

```
print(X)
```

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak	\
0	63	1	3	145	233	1	0	150	0	2.3	
1	37	1	2	130	250	0	1	187	0	3.5	
2	41	0	1	130	204	0	0	172	0	1.4	

 ×

 Your download's being scanned.  
We'll let you know if there's an issue.

```

3      56      1      1      120      236      0      1      178      0      0.8
4      57      0      0      120      354      0      1      163      1      0.6
..      ...      ...      ..      ...      ...      ...      ...      ...      ...
298    57      0      0      140      241      0      1      123      1      0.2
299    45      1      3      110      264      0      1      132      0      1.2
300    68      1      0      144      193      1      1      141      0      3.4
301    57      1      0      130      131      0      1      115      1      1.2
302    57      0      1      130      236      0      0      174      0      0.0

```

```

      slope  ca  thal
0         0   0    1
1         0   0    2
2         2   0    2
3         2   0    2
4         2   0    2
..      ...  ..   ...
298      1   0    3
299      1   0    3
300      1   2    3
301      1   1    3
302      1   1    2

```

```
[303 rows x 13 columns]
```

```
print(Y)
```

```

0      1
1      1
2      1
3      1
4      1
..
298    0
299    0
300    0
301    0
302    0
Name: target, Length: 303, dtype: int64

```

### Splitting the Data into Training data & Test Data

```
X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size=0.2, stratify=Y, random_state=2)
```

```
print(X.shape, X_train.shape, X_test.shape)
```

```
(303, 13) (242, 13) (61, 13)
```

### Model Training

#### Logistic Regression

```
model = LogisticRegression()
```

```
# training the LogisticRegression model with Training data
model.fit(X_train, Y_train)
```

```

/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_logistic.py:469: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

```

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

```
n_iter_i = _check_optimize_result(
```

```
    LogisticRegression
```

```
LogisticRegression()
```

### Model Evaluation

#### Accuracy Score



McAfee | WebAdvisor

Your download's being scanned.  
We'll let you know if there's an issue.

```
# accuracy on training data
X_train_prediction = model.predict(X_train)
training_data_accuracy = accuracy_score(X_train_prediction, Y_train)
```

```
print('Accuracy on Training data : ', training_data_accuracy)
```

```
↗ Accuracy on Training data : 0.8512396694214877
```

```
# accuracy on test data
X_test_prediction = model.predict(X_test)
test_data_accuracy = accuracy_score(X_test_prediction, Y_test)
```

```
print('Accuracy on Test data : ', test_data_accuracy)
```

```
↗ Accuracy on Test data : 0.819672131147541
```

### Building a Predictive System

```
input_data = (62,0,0,140,268,0,0,160,0,3.6,0,2,2)
```

```
# change the input data to a numpy array
input_data_as_numpy_array= np.asarray(input_data)
```

```
# reshape the numpy array as we are predicting for only on instance
input_data_resaped = input_data_as_numpy_array.reshape(1,-1)
```

```
prediction = model.predict(input_data_resaped)
print(prediction)
```

```
if (prediction[0]== 0):
    print('The Person does not have a Heart Disease')
else:
    print('The Person has Heart Disease')
```

```
↗ [0]
The Person does not have a Heart Disease
/usr/local/lib/python3.10/dist-packages/sklearn/base.py:493: UserWarning: X does not have valid feature names, but LogisticRegression wa
warnings.warn(
```



McAfee | WebAdvisor



Your download's being scanned.  
We'll let you know if there's an issue.